

# Lecture 7

## Naive Bayes and Sentiment Classification



**CS 6320**

# Naive Bayes and Sentiment Classification

---

- **Classification** is a basic activity of human life.  
Recognize and classify. Examples: classify people, images, voices, assigning grades, emails, etc.
- **Text classification/categorization** – is the task of assigning a label (category) to a document ; it is a basic NLP capability.
- **Sentiment analysis**- is an example of text categorization – it assigns a positive or negative label to text.  
Examples: Product reviews, movie review, book review, restaurant reviews, political commentaries, etc.

## Other text classification problems

- Spam detection
- Language identification

# Naive Bayes and Sentiment Classification

---

- Example

- + *...richly applied satire, and some great plot twists*
- *It was pathetic. The worst part about it was the boxing scenes...*
- + *...awesome caramel sauce and sweet toasty almonds. I love this place!*
- *...awful pizza and ridiculously overpriced...*

# Approach to classification: supervised machine learning

---

- Take an input  $x$  and a fixed set of output classes  $Y = y_1, y_2, \dots, y_M$  and return a predicted class  $y \in Y$ .

For texts:  $c$  is a class, ( instead of  $y$  ) and  $d$  – document.

Training examples  $(d_1, c_1), (d_2, c_2) \dots (d_N, c_N)$

Goal: classify a new document  $d$  to its correct class  $c \in C$ .

# Probabilistic classifiers

---

1. **Generation classifiers** – given an observation return the class most likely to have generated the observation.

Example: Naive Bayes

2. **Discriminative classifiers** – learn what features from input are most useful to discriminate between possible classes.

Example: logistic regression

# Naive Bayes Classifiers

---

- Given a document  $d$  and set of classes  $c \in C$ , returns

$$\hat{c} = \operatorname{argmax}_{c \in C} P(c|d)$$

$$= \operatorname{argmax}_{c \in C} \frac{P(d|c) P(c)}{P(d)}$$

$$= \operatorname{argmax}_{c \in C} P(d|c) P(c)$$

$\uparrow$                        $\uparrow$   
 likelihood          prior  
 probability        probability

# Naive Bayes Classifiers

---

- A document  $d$  is represented by a set of features  $f_1, f_2, \dots, f_n$

$$\hat{c} = \operatorname{argmax} P(f_1, f_2, \dots, f_n | c) P(c)$$

- Naive Bayes assumption: Features are independent, ie:  $P(f_i | c)$  are independent.

$$P(f_1, f_2, \dots, f_n | c) = P(f_1 | c) P(f_2 | c) \dots P(f_n | c)$$

$$c_{NB} = \operatorname{argmax}_{c \in C} P(c) \prod_{f \in F} P(f | c)$$

# Naive Bayes Classifiers

---

- Bag -of-words approach: Features use words.
- Words are specified by their positions in the document;  
 $w_i$ — word at position  $i$

$$C_{NB} = \operatorname{argmax}_{c \in C} P(c) \prod_{i \in \text{positions}} P(w_i | c)$$

- Use log space to avoid underflow and increase speed.

$$C_{NB} = \operatorname{argmax}_{c \in C} [\log P(c) + \sum_{i \in \text{positions}} \log P(w_i | c)]$$



# Training the Naïve Bayes Classifiers

---

- How to learn  $P(c)$  and  $P(f_i|c)$ ?

$N_c$  – percentage of documents in training set with class  $c$ .

$N_{doc}$  - total number of documents.

$$\hat{P}(c) = \frac{N_c}{N_{doc}}$$

- $$\hat{P}(w_i|c) = \frac{\text{count}(w_i, c)}{\sum_{w \in V} \text{count}(w, c)}$$

$V$  - vocabulary for all classes.

Example: 
$$\hat{P}(\text{fantastic}|+) = \frac{\text{count}(\text{fantastic}, +)}{\sum_{w \in V} \text{count}(w, +)}$$

- Use Laplace smoothing

$$\hat{P}(w_i|c) = \frac{\text{count}(w_i, c) + 1}{\sum_{w \in V} \text{count}(w, c) + V}$$

# Example

	Cat	Documents
Training	-	just plain boring
	-	entirely predictable and lacks energy
	-	no surprises and very few laughs
	+	very powerful
	+	the most fun film of the summer
Test	?	S=predictable with no fun

$$P(-) = \frac{3}{5}$$

$$P(+) = \frac{2}{5}$$

$$P(\text{"predictable"}|-) = \frac{1+1}{14+20} \quad P(\text{"predictable"}|+) = \frac{0+1}{9+20}$$

$$P(\text{"no"}|-) = \frac{1+1}{14+20} \quad P(\text{"no"}|+) = \frac{0+1}{9+20}$$

$$P(\text{"fun"}|-) = \frac{0+1}{14+20} \quad P(\text{"fun"}|+) = \frac{1+1}{9+20}$$

$$P(-)P(S|-) = \frac{3}{5} \times \frac{2 \times 2 \times 1}{34^3} = 6.1 \times 10^{-5}$$

$$P(+ )P(S|+) = \frac{2}{5} \times \frac{1 \times 1 \times 2}{29^3} = 3.2 \times 10^{-5}$$

→ Pick Negative class

# Improvements

---

- Fix negations

*didn't like this movie*

*didn't* becomes     *Not\_like,*  
                              *Not\_this,*  
                              *Not\_movie w*

- Sentiment lexicons – list of words that are pre-annotated with positive and negative sentiment.  
Examples: MPQA subjectivity lexicon; 6885 words, 2718 positive and 4912 negative; with strong or weak bias.
  - + *admirable, beautiful, confident, dazzling, great...*
  - *awful, bad, cheat, deny, envious, foul, hate...*

# Naive Bayes as a Language Model

- A language model predicts the next word after seeing a string of words; and calculates probability of strings.
- Naïve Bayes assigns a probability to each word  $P(\text{word}|c)$ . It can also assign a probability to each sentence

$$P(s | c) = \prod_{i \in \text{positions}} P(w_i | c)$$

	$P(w +)$	$P(w -)$
I	0.1	0.2
love	0.1	0.001
this	0.01	0.01
fun	0.05	0.005
film	0.1	0.1

$$P("I \text{ love this fun film"}|+) = 0.1 \times 0.1 \times 0.01 \times 0.05 \times 0.1 = 0.0000005$$

$$P("I \text{ love this fun film"}|-) = 0.2 \times 0.001 \times 0.01 \times 0.005 \times 0.1 = 0.000000001$$

$$P(s|+) > P(s|-)$$

# Evaluation: Precision and Recall

		<i>gold standard labels</i>		
		gold positive	gold negative	
<i>system output labels</i>	system positive	true positive	false positive	$precision = \frac{tp}{tp + fp}$
	system negative	false negative	true negative	
		$recall = \frac{tp}{tp + fn}$		$accuracy = \frac{tp + tn}{tp + fp + tn + fn}$

$$Precision = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

$$Recall = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

F – measure

$$F_{\beta} = \frac{(\beta^2 + 1)PR}{\beta^2 P + R}$$

$\beta > 1$  favors recall

$\beta < 1$  favors precision

# More than two classes

---

- Two Possibilities:

1. Multi-label classification (*any-of*)

- Each document can be assigned more than one label
- Build separate binary classifiers for each class  $c$ , train on positive examples labeled  $c$  and negative examples not labeled  $c$ .
- For a test document  $d$ , each classifier makes its decision independently, and multiple labels can be assigned to  $d$ .

# More than two classes

---

## 2. Multinomial classification (*one-of*)

- Classes are mutually exclusive
- Build separate binary classifiers
- For a test document, run all classifiers and choose the label from the classifier with the highest score.

# Example: 3-way one-of email classification

		<i>gold labels</i>			
		<b>Urgent</b>	<b>Normal</b>	<b>Spam</b>	
<i>System output</i>	urgent	8	10	1	$precision_u = \frac{8}{8 + 10 + 1}$
	normal	5	60	50	$precision_n = \frac{60}{5 + 60 + 50}$
	spam	3	30	200	$precision_s = \frac{200}{3 + 30 + 200}$
		$recall_u = \frac{8}{8 + 5 + 3}$	$recall_n = \frac{60}{10 + 60 + 30}$	$recall_s = \frac{200}{1 + 50 + 200}$	



# Example: 3-way one-of email classification

	Class 1: Urgent		Class 2: Normal		Class 3: Spam		Pooled	
	true urgent	true not	true normal	true not	true spam	true not	true yes	true no
System urgent	8	11	60	55	200	33	268	99
System not	8	340	40	212	51	83	99	635
$precision = \frac{8}{8 + 11} = .42$								
$precision = \frac{60}{60 + 55} = .52$								
$precision = \frac{200}{200 + 33} = .86$								
$macroaverage\ precision = \frac{.42 + .52 + .86}{3} = .60$								
$microaverage\ precision = \frac{268}{268 + 99} = .73$								