

CS 6375 Machine Learning

Professor: Sriraam Natarajan

Email: Sriraam.Natarajan@utdallas.edu

Webpage:

<http://utdallas.edu/~sxn177430/Courses/6375ML.html>

Class Hours: MW 11:30-12:45

Office Hours W 10:30-11:30 and by appointment

Course Information

- No text book required, slides and reading materials will be provided in elearning page
- There are a few recommended books that are good references
 - Pattern recognition and machine learning by Chris Bishop (Bishop)
 - Machine learning by Tom Mitchell (TM)
 - Machine Learning: A Probabilistic Perspective by Kevin Murphy (Murphy)
- Slides will be posted **after** the classes in time for assignments
- Most of the slides are based on Tom Dietterich and Jude Shavlik's class slides (obtained with their permission)

Course Syllabus

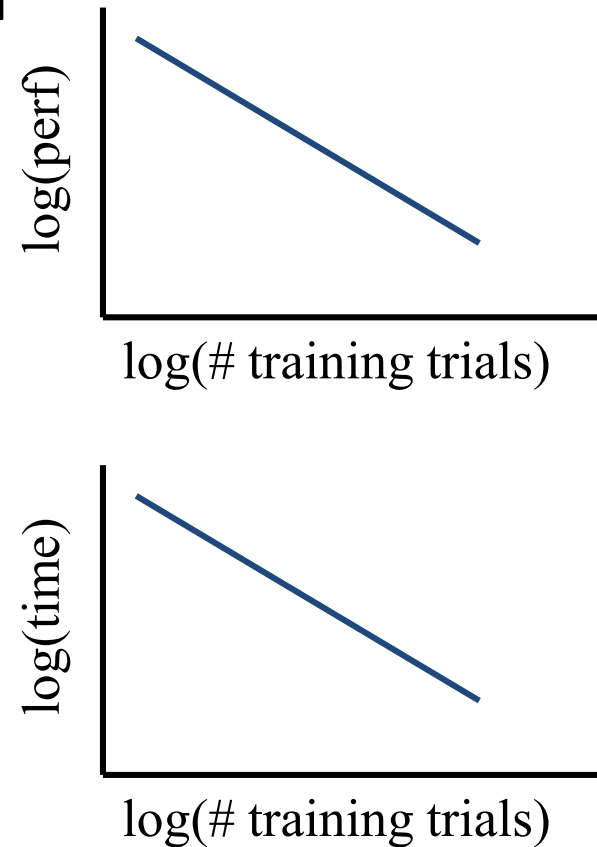
- Linear Classifiers (Naive Bayes, and logistic regression)
- Non-linear classifiers (neural networks, decision trees, support-vector machines, nearest neighbor methods)
- Ensemble Methods (bagging and boosting)
- Computational Learning Theory
- Reinforcement Learning
- Unsupervised Learning

Grading

- Assignments (25%)
- Programming Assignments(25%)
- Midterm (25%)
- Final (25%)

Learning

- **Herbert Simon:** “Learning is any process by which a system improves performance from experience.”
- **A Collaborator:** “We need machine learning because we like being lazy. i.e., let the machines learn to do what we do”
- **Ray Mooney:** Develop systems that are too difficult/expensive to construct manually because they require specific detailed skills or knowledge tuned to a specific task.
- Develop systems that can automatically adapt and customize themselves to individual users.
- Discover new knowledge from large databases.



History of Machine Learning

- 1950s
 - Samuel's checker player
 - Selfridge's Pandemonium
- 1960s:
 - Neural networks: Perceptron
 - Pattern recognition
 - Learning in the limit theory
 - Minsky and Papert prove limitations of Perceptron
- 1970s:
 - Symbolic concept induction
 - Winston's arch learner
 - Expert systems and the knowledge acquisition bottleneck
 - Quinlan's ID3
 - Michalski's AQ and soybean diagnosis
 - Scientific discovery with BACON
 - Mathematical discovery with AM

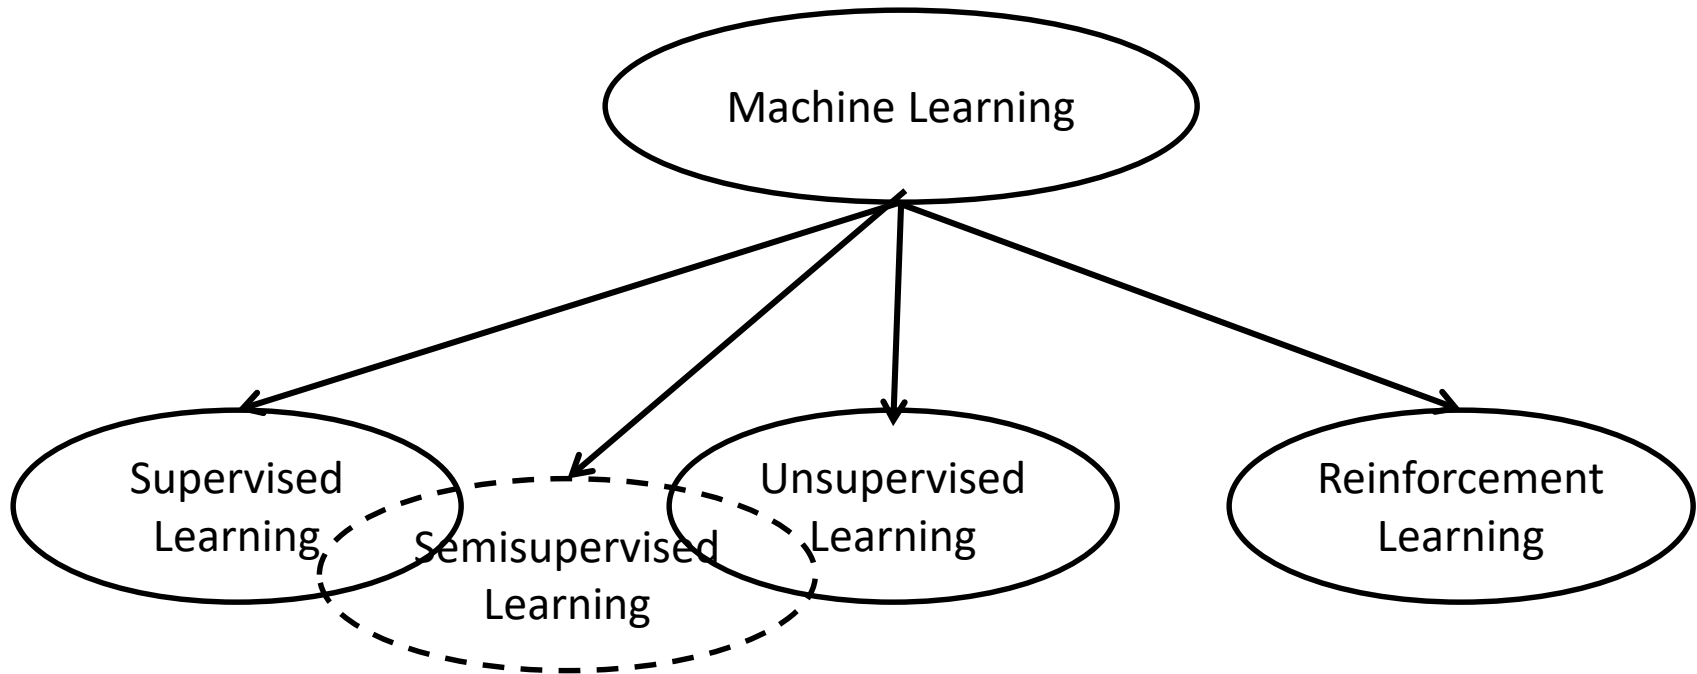
History of Machine Learning (cont.)

- 1980s:
 - Advanced decision tree and rule learning
 - Explanation-based Learning (EBL)
 - Learning and planning and problem solving
 - Utility problem
 - Analogy
 - Cognitive architectures
 - Resurgence of neural networks (connectionism, backpropagation)
 - Valiant's PAC Learning Theory
 - Focus on experimental methodology
- 1990s
 - Data mining
 - Adaptive software agents and web applications
 - Text learning
 - Reinforcement learning (RL)
 - Inductive Logic Programming (ILP)
 - Ensembles: Bagging, Boosting, and Stacking
 - Bayes Net learning

History of Machine Learning (cont.)

- 2000s
 - Support vector machines
 - Kernel methods
 - Graphical models
 - Statistical relational learning
 - Transfer learning
 - Sequence labeling
 - Collective classification and structured outputs
 - Computer Systems Applications
 - Compilers
 - Debugging
 - Graphics
 - Security (intrusion, virus, and worm detection)
 - E mail management
 - Personalized assistants that learn
 - Learning in robotics and vision

Machine Learning



Supervised Learning – Medical Diagnosis

Do Not Have Diabetes

blood glucose = 30

body mass index = 120 kg/m²

diastolic bp = 79 mm Hg

age = 32 years



blood glucose = 22

body mass index = 160 kg/m²

diastolic bp = 80 mm Hg

age = 63 years



blood glucose = 22

body mass index = 160 kg/m²

diastolic bp = 80 mm Hg

age = 18 years



blood glucose = 40

body mass index = 160 kg/m²

diastolic bp = 80 mm Hg

age = 63 years



blood glucose = 30
body mass index = 120 kg/m²
diastolic bp = 73 mm Hg
age = 27 years



blood glucose = 46

body mass index = 150 kg/m²

diastolic bp = 110 mm Hg

age = 55 years



blood glucose = 45
body mass index = 180 kg/m²

diastolic bp = 95 mm Hg

age = 49 years



blood glucose = 21


body mass index = 140 kg/m²

diastolic bp = 99 mm Hg

age = 37 years


Have Diabetes

Supervised Learning – Digit Recognition




Unsupervised Learning


- Also, called clustering




Unsupervised Learning – Object Recognition



Reinforcement Learning




Reinforcement Learning – Robocup Soccer



Supervised Learning


- **Given**: Training examples $\langle x, f(x) \rangle$ for some unknown function f .
- **Find**: A good approximation to f .
- Situations where there is no human expert
 - x : bond graph of a new molecule
 - $f(x)$: predicted binding strength to AIDS protease molecule
- Situations where humans can perform the task but can't describe how they do it
 - x : picture of a hand-written character
 - $f(x)$: ascii code of the character
- Situations where the desired function is changing frequently
 - x : description of stock prices and trades for last 10 days
 - $f(x)$: recommended stock transactions
- Situations where each user needs a customized function f
 - x : incoming email message
 - $f(x)$: importance score for presenting to the user (or deleting without presenting)

Supervised Learning



- Training examples: drawn independently at random according to unknown distribution $P(x, y)$
- The learning algorithm analyzes the training examples and produces a classifier f
- Given a new point $\langle x, y \rangle$ drawn from P , the classifier is given x and predicts $\hat{y} = f(x)$
- The loss $L(\hat{y}, y)$ is then measured
- Goal of the learning algorithm: Find the f that minimizes the expected loss $E_{P(x, y)}[L(f(x), y)]$

Example – Curve fitting



The underlying function:

$$t = \sin(2\pi x) + \varepsilon$$

where ε is Gaussian noise

- Input: The Blue circles
 - The circles are generated using the green curve
- Goal: Make **accurate** predictions on new **unseen** points such that some notion of error is minimized
 - Predict values of t given values of x

Curve fitting - Continued

- There are ***infinitely*** many possible functions that can fit this data.
- Hence, we need to focus on a smaller number of functions possible
 - This is called **hypothesis** space
- Ex., we can restrict the order of polynomial (nth order)


$$f(x, w) = w_0 + w_1x + w_2x^2 + \dots + w_nx^n$$

- We wish to learn the parameters $\langle w_0, w_1, \dots, w_n \rangle$
- These parameters must be learned such that some loss function is minimized

- For example, squared error
$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(x_n, \mathbf{w}) - t_n)^2$$

Let us not worry about how it is solved yet ☺ These can be solved easily by optimization techniques.


Which Model to choose?



- Red curve is the one with different values for n
- Which n should we choose – Model selection
- Can we use the error $E(\mathbf{w})$ as the criterion?

Overfitting

- As n increases, training error decreases monotonically
- As n increases test error can increase
- Test error can decrease at first, but increases
- Overfitting can occur
 - When the model is too complex and trivially fits the data (i.e., too many parameters)
 - When the data is not enough to estimate the parameters
 - Model captures the noise (or the chance)



Terminology

Features

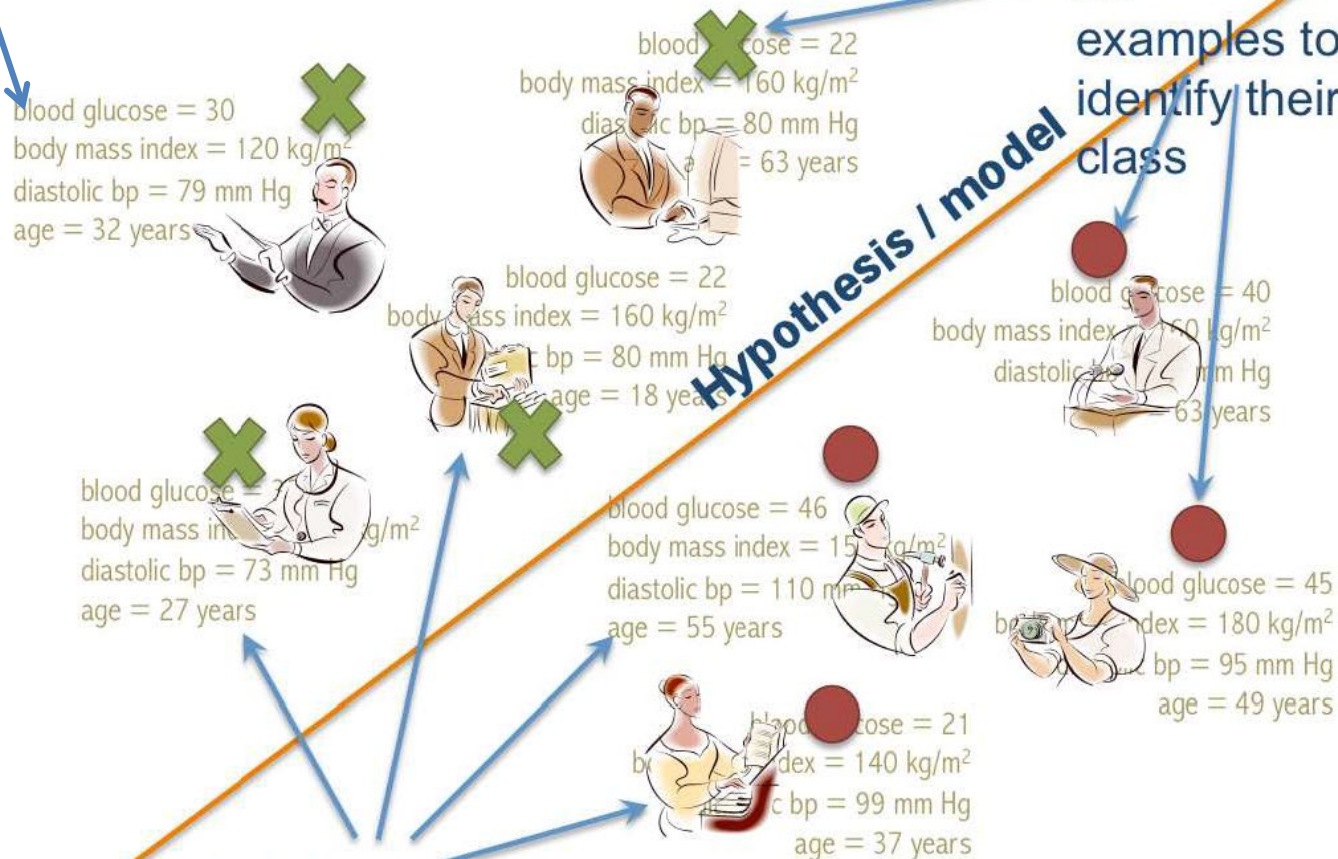
Do Not Have Diabetes

training labels
for
examples to
identify their
class

Hypothesis / model

training examples

Have Diabetes



Terminology

- **Training Example:** $\langle \mathbf{x}, y \rangle$
 - \mathbf{x} : feature vector (describes the attributes of something)
 - y : label (continuous values for regression problems: $[1, 2, \dots, k]$ for classification problems)
- **Training set** A set of training examples drawn randomly from $P(\mathbf{x}, y)$
 - **Key Assumption:** Independent and identically distributed. i.e., all the examples are drawn from the same distribution but are drawn independent of one another
- **Target function** True mapping from \mathbf{x} to y
- **Hypothesis:** A function h considered by the learning algorithm to be similar to the target function
- **Test set:** A set of examples drawn from $P(\mathbf{x}, y)$ to evaluate the “goodness of h ”
- **Hypothesis Space:** The space of all hypotheses that can in principle be considered and returned by the learning algorithm

Key approaches

- Directly learn a mapping $y = f(\mathbf{x})$
 - No uncertainty is captured
- Learn the joint distribution i.e., learn $p(y, \mathbf{x})$
 - Captures uncertainty about both the attributes \mathbf{x} and the target y
- Learn the conditional distribution i.e., learn $p(y | \mathbf{x})$
 - $p(\mathbf{x}, y) = p(y | \mathbf{x})p(\mathbf{x})$
 - Hence this avoids modeling the distribution of \mathbf{x}
 - In general, this is akin to assuming a uniform distribution over \mathbf{x}
 - Can also be considered as saying “I do not care about \mathbf{x} but only $P(y | \mathbf{x})$ ”
- Once we learn p , how do we choose y ? This is called as decision-theory

Example Training data

x0	x1	x2	x3	y
0	0	1	1	0
1	0	1	0	1
1	1	0	0	0
0	0	1	1	0
0	1	0	0	1
1	0	1	1	1
1	1	0	1	0

Homework Q: Is there a function that is a conjunction of literals that fits the data perfectly?

Two Views

- View1: Learning is the removal of our remaining uncertainty
 - Suppose we know that the function is a conjunction. Then we could use the training example to determine what this function is
- View2: Learning requires guessing a good small hypothesis class
 - We could start with a small hypothesis class and enlarge it till it contains an hypothesis that fits the data
 - For instance, we could say it is a conjunction of 2 literals, then expand to 3 and then say it is may be a disjunction of conjunctions and so on
- But the problem is that we could be wrong
 - Our prior knowledge might be wrong
 - Our guess of hypothesis class itself is wrong – smaller the class higher is the chance of being wrong

Key Questions

- What are good hypothesis spaces?
- What algorithms work on these spaces?
- How can we generalize to unseen points (i.e., avoid overfitting)?
- How can we trust our results?
- Are some problems computationally intractable?
- How can we formulate practical problems as Machine learning ones?

Learning Algorithms

- Search
 - Direct Computation
 - Local Search: Start with an initial hypothesis, make gradual improvements till a local maximum is reached
 - Constructive search: start with an empty hypothesis and grow it till local maximum
- Timing
 - Eager: Analyze training data and construct hypothesis
 - Lazy: Store the training data and wait till a test point is presented, then construct an hypothesis to classify the test point
- Online vs Batch (for eager)
 - Online: Analyze each training point as they arrive
 - Batch: Collect all examples, analyze them together

Next

- Learn a classifier (directly the function f)
 - Learn a conditional distribution: Logistic Regression
 - Learn the joint distribution: Linear discriminant analysis
-
- **Key question: How many of you know some basics of probability, maximum likelihood, density estimation etc?**