

# Linear Threshold Units

# Key approaches

- Directly learn a mapping  $y = f(\mathbf{x})$ 
  - No uncertainty is captured
- Learn the joint distribution i.e., learn  $p(y, \mathbf{x})$ 
  - Captures uncertainty about both the attributes  $\mathbf{x}$  and the target  $y$
- Learn the conditional distribution i.e., learn  $p(y | \mathbf{x})$ 
  - $p(\mathbf{x}, y) = p(y | \mathbf{x})p(\mathbf{x})$
  - Hence this avoids modeling the distribution of  $\mathbf{x}$
  - In general, this is akin to assuming an uniform distribution over  $\mathbf{x}$
  - Can also be considered as saying “I do not care about  $\mathbf{x}$  but only  $P(y | \mathbf{x})$ ”
- Once we learn  $p$ , how do we choose  $y$ ? This is called as decision-theory

# Linear Threshold Units

$$h(x) = \begin{cases} +1 & \text{if } w_1x_1 + \dots + w_nx_n \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

Current Assumption: Each feature  $x_i$  and each weight  $w_j$  is a real number

Three Algorithms:

1. Perceptron – Directly learns the function
2. Logistic Regression: Conditional Distribution
3. Linear Discriminant Analysis: Joint Distribution

# What can an LTU represent

- Conjunctions

$$x_1 \wedge x_2 \wedge x_4 \Leftrightarrow y$$

$$1 \cdot x_1 + 1 \cdot x_2 + 0 \cdot x_3 + 1 \cdot x_4 \geq 3$$

- At least m-of-n

$$\text{at-least-2-of } \{x_1, x_3, x_4\} \Leftrightarrow y$$

$$1 \cdot x_1 + 0 \cdot x_2 + 1 \cdot x_3 + 1 \cdot x_4 \geq 2$$

# What cannot be represented

- Complex disjunctions

$$(x_1 \wedge x_2) \vee (x_3 \wedge x_4) \Leftrightarrow y$$

- Exclusive-OR

$$(x_1 \wedge \neg x_2) \vee (\neg x_1 \wedge x_2) \Leftrightarrow y$$

# A Canonical Representation

- Given a training example:  $(\langle x_1, x_2, x_3, x_4 \rangle, y)$   $y \in \{-1, 1\}$
- Transform it to canonical representation

$$(\langle 1, x_1, x_2, x_3, x_4 \rangle, y)$$

- Learn a linear function  $g(\mathbf{x}, \mathbf{w}) = \mathbf{w}^T \mathbf{x}$ , where  $\mathbf{w} = \langle w_0, w_1, w_2, w_3, w_4 \rangle$
- Each  $\mathbf{w}$  corresponds to one hypothesis

$$h(\mathbf{x}) = \text{sign}(g(\mathbf{x}, \mathbf{w}))$$

- A prediction is correct if  $y \mathbf{w}^T \mathbf{x} > 0$
- Goal of learning is to find a good  $\mathbf{w}$ 
  - e.g., a  $\mathbf{w}$  such that  $h(\mathbf{x})$  makes few mis-predictions

# LTU Hypotheses space


- Fixed size
- Deterministic
- Continuous parameters

# Geometric View

- Let us consider 3 training examples:

$\langle 1.0, 1.0 \rangle, +1$   $\langle 0.5, 3.0 \rangle, +1$   $\langle 2.0, 2.0 \rangle, -1$

The classifier should look like the following






# The discriminant function is a hyperplane

- The equation  $u(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x}$  is a plane

$$\hat{y} = \begin{cases} +1 & \text{if } u(x) \geq 0 \\ -1 & \text{otherwise} \end{cases}$$



# We can view this problem as an optimization problem


- Formulate learning problem as an optimization problems
  - Given:
    - A set of  $N$  training examples  $\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
    - A loss function  $L$
  - Find the weight vector  $\mathbf{w}$  that minimizes the objective function - the expected/average loss on training data

$$J(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N L(\mathbf{w} \cdot \mathbf{x}_i, y_i)$$

- Many machine learning algorithms apply some optimization algorithm to find a good hypothesis.

# 0/1 Loss function

- 0/1 Loss function:  $J_{0/1}(w) = \frac{1}{N} \sum_{i=1}^N L(\text{sgn}(w \cdot x_i), y_i)$   
 $L(y', y) = 0$  when  $y' = y$ , otherwise  $L(y', y) = 1$
- Does not produce useful gradient since the surface of  $J$  is flat




# Other alternative – modified hinge loss


- Instead we will consider the “**perceptron criterion**” (a slightly modified version of hinge loss):

$$J_p(w) = \frac{1}{N} \sum_{i=1}^N \max(0, -y_i w \cdot x_i)$$

- The term  $\max(0, -y_i w \cdot x_i)$  is 0 when  $y_i$  is predicted correctly otherwise it is equal to the “confidence” in the mis-prediction
- Has a nice gradient leading to the solution region




0/1 loss



Perceptron  
criterion

# Gradient Descent minimizes the loss function



- Start with weight vector  $\mathbf{w} = (w_0, \dots, w_n)$
- Compute gradient  $\nabla J(\mathbf{w}_0) = \left( \frac{\partial J(\mathbf{w})}{\partial w_0}, \frac{\partial J(\mathbf{w})}{\partial w_1}, \dots, \frac{\partial J(\mathbf{w})}{\partial w_n} \right)_{\mathbf{w}_0}$
- Compute  $w_1 = w_0 - \eta \nabla J(w_0)$  where  $\eta$  is "step size"
- Repeat until convergence

# Gradient Descent

- The objective function consists of a sum over data points--- we can update the parameter after observing each example
- This is referred to as Stochastic gradient descent approach

$$J(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \max(0, -y_i \mathbf{w} \cdot \mathbf{x}_i)$$

$$J_i(\mathbf{w}) = \max(0, -y_i \mathbf{w} \cdot \mathbf{x}_i)$$

$$\frac{\partial J_i}{\partial w_j} = \begin{cases} 0 & \text{if } y_i \mathbf{w} \cdot \mathbf{x}_i > 0 \\ -y_i x_{ij} & \text{otherwise} \end{cases}$$

$$\nabla J_i = \begin{cases} 0 & \text{if } y_i \mathbf{w} \cdot \mathbf{x}_i > 0 \\ -y_i \mathbf{x}_i & \text{otherwise} \end{cases}$$

After observing  $(\mathbf{x}_i, y_i)$ , if it is a mistake  $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$

# Online Perceptron Algo

Let  $\mathbf{w} = (0,0,0,...,0)$  be the initial weight vector

Repeat forever

Accept training example  $i$ :  $\langle \mathbf{x}_i, y_i \rangle$

$$u_i = \mathbf{w} \cdot \mathbf{x}_i$$

IF  $(y_i \cdot u_i < 0)$

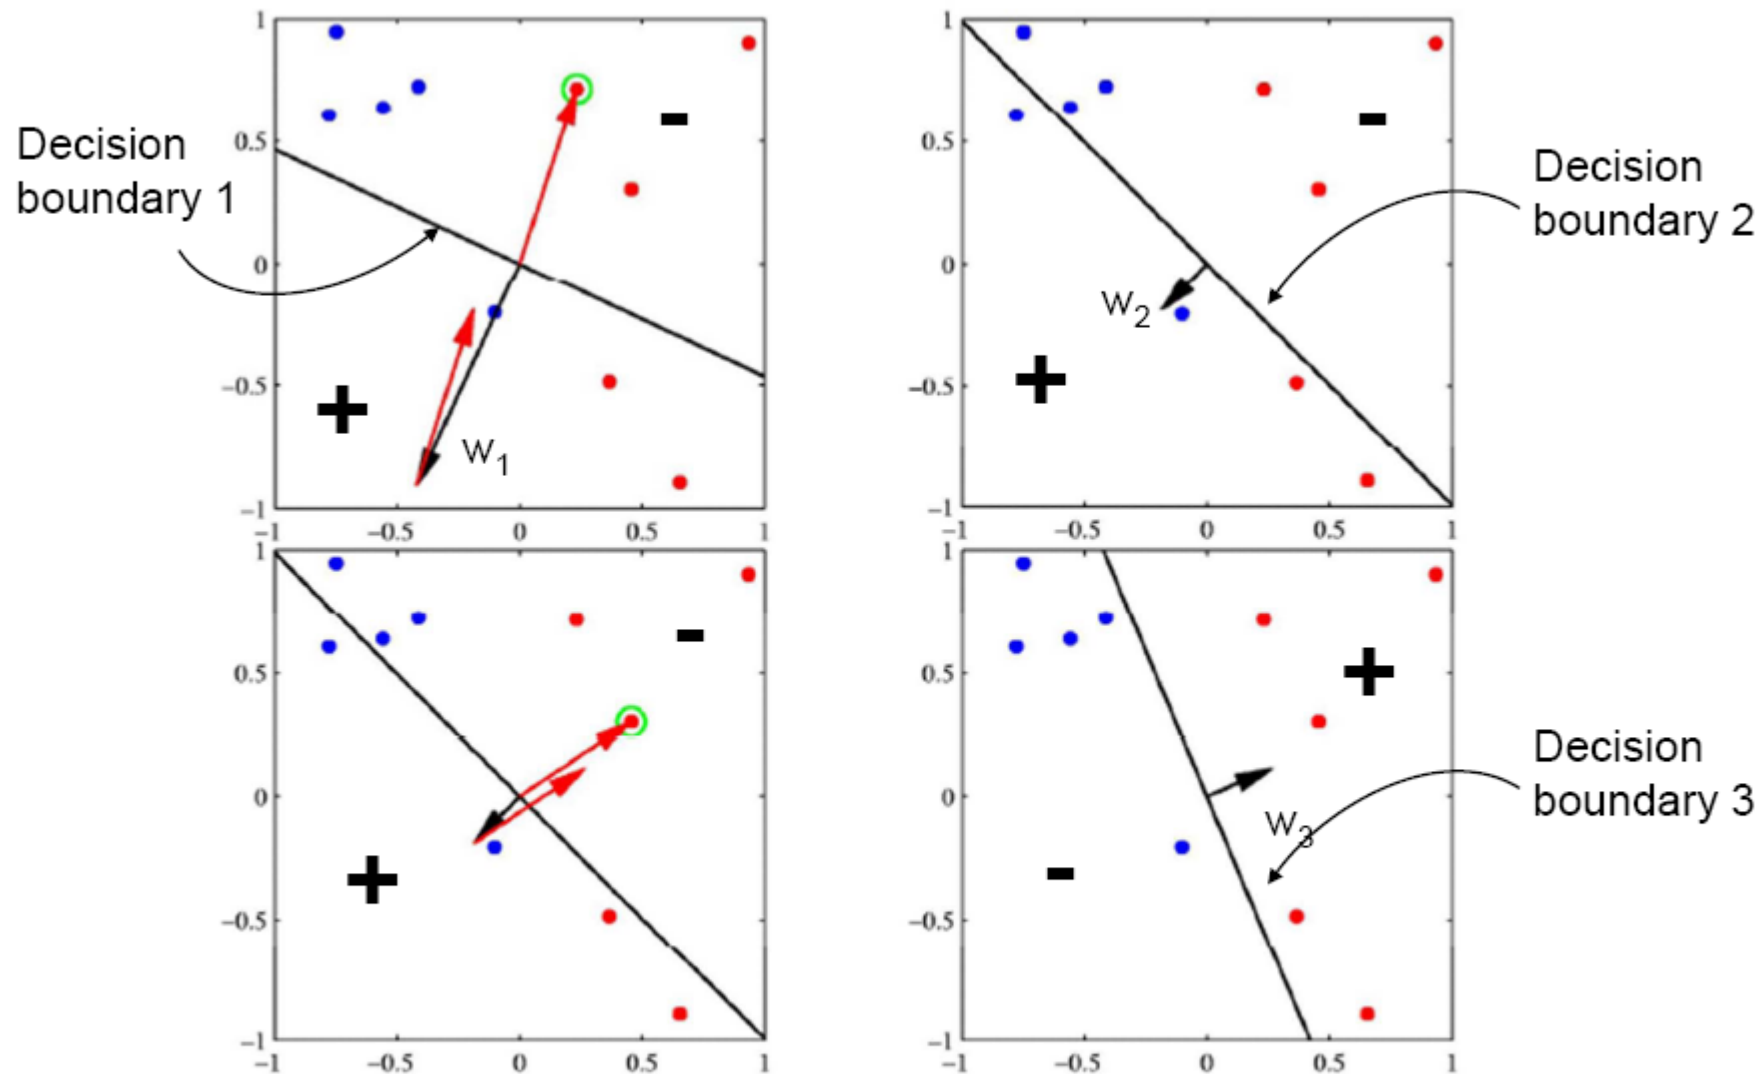
For  $j = 1$  to  $n$  do // For every feature, compute gradient

$$g_j := y_i \cdot x_{ij}$$

$$\mathbf{w} := \mathbf{w} + \eta \mathbf{g}$$

This is called stochastic gradient descent as the overall gradient is approximated by the gradient from each example

When an error is made, moves the weight in a direction that corrects the error



Red points belong to the positive class,  
blue points belong to the negative class



# Batch Perceptron Algo

Let  $w = (0,0,0,...,0)$  be the initial weight vector

Let  $g = (0,0,0,...,0)$  be the initial gradient vector

Repeat until convergence

For  $i = 1$  to  $N$  do

$$u_i = \mathbf{w} \cdot \mathbf{x}_i$$

IF  $(y_i \cdot u_i < 0)$

For  $j = 1$  to  $n$  do

$$g_j := g_j - y_i \cdot x_{ij}$$

$$g := g / N$$

$$w := w - \eta g$$

When  $\eta = 1$  it is a fixed increment perceptron

# Step size

- Referred to as learning rate -- an important factor in many learning algorithms
- Learning rate must decrease to zero in order for the algorithm to converge

$$\lim_{t \rightarrow \infty} \eta_t = 0$$


$$\sum_{t=0}^{\infty} \eta_t = \infty$$

$$\sum_{t=0}^{\infty} \eta_t^2 < \infty$$

- Some optimization algorithms set the step size automatically and converge faster
- For LTUs, there is only one basin. i.e., local minimum is global minimum. Choosing good step size will result in faster convergence

# Decision Boundaries


- A classifier can be viewed as partitioning the input space or feature  $X$  into decision regions



- A set of points that can be separated by a linear decision boundary is said to be linearly separable.

# Not Linearly Separable

- X-Or



## Perceptron – Key Result

If training set is linearly separable, perceptron WILL find it in finite steps

If training data is not linearly separable, perceptron will never converge


# Summary of Perceptron

- Directly learns a classifier
- Local Search
  - Begins with an initial weight vector. Modifies it iteratively to minimize a loss function. The error function is loosely related to the goal of minimizing classification errors
- Eager
  - The classifier is constructed from training examples
  - The examples can then be discarded
- Online or Batch versions


# Limitations of Linear Hypotheses

**Regression:** Non-linear regression functions

**Classification:** Linearly inseparable classes



Explicitly transform the data  $x \rightarrow (x, x^2)$





Explicitly transform the data  $(x_1, x_2) \rightarrow (x_1^2, x_2^2, \sqrt{2}x_1x_2)$

**A potential solution:**  
Transforming the data into a **higher-dimensional** space makes the problem linear in that space at expense of adding more features.

Coming up with such transformations is not easy.

More on this later in the course!



# Next

- Logistic Regression
- Linear Discriminant Analysis
- **First Assignment:** Due the following week